

20 Advanced Topics

1. Normalization v/s Denormalization

Normalization

- Normalization is a process of organizing data in a database to reduce redundancy and improve data integrity
- It involves dividing a database into two or more tables and defining relationships between the tables
- The main goals of normalization are to eliminate redundant data, ensure data dependencies make sense, and reduce the potential for data anomalies

Example

- Consider a university database with information about students, courses, and enrollments

Before Normalization

StudentID	StudentName	CourseID	CourseName	Professor
1	Alice	C101	Math	Dr. Smith
2	Bob	C102	Physics	Dr. Johnson
1	Alice	C103	Chemistry	Dr. Lee

After Normalization (3NF)

Students Table

StudentID	StudentName
1	Alice
2	Bob

Courses Table

CourseID	CourseName	Professor
C101	Math	Dr. Smith
C102	Physics	Dr. Johnson
C103	Chemistry	Dr. Lee

Enrollments Table

StudentID	CourseID
1	C101
2	C102
1	C103

Denormalization

- Denormalization is the process of combining normalized tables into a single table to improve read performance
- This can lead to redundant data and potential anomalies but can significantly speed up data retrieval

Example

- Consider the same university database
- To improve read performance for a specific query, we might combine the Students, Courses, and Enrollments tables back into a single table

After Denormalization

StudentID	StudentName	CourseID	CourseName	Professor
1	Alice	C101	Math	Dr. Smith
2	Bob	C102	Physics	Dr. Johnson
1	Alice	C103	Chemistry	Dr. Lee

2. Advanced Query Optimization

- Advanced query optimization involves techniques to improve the performance of SQL queries
- This can include indexing, query rewriting, and using execution plans

Example

- Suppose we have a table `orders` with millions of rows, and we frequently query orders by `customer_id`

Initial Query

```
SELECT * FROM orders WHERE customer_id = 12345;
```

Optimization Techniques

- **Indexing:** Create an index on the `customer_id` column

```
CREATE INDEX idx_customer_id ON orders (customer_id);
```

- **Query Rewriting:** Simplify complex queries to make them more efficient

```
SELECT order_id, order_date FROM orders WHERE customer_id = 12345;
```

- **Execution Plan:** Analyze the query execution plan to identify bottlenecks

```
EXPLAIN SELECT * FROM orders WHERE customer_id = 12345;
```

3. Partitioning Tables

- Partitioning is a database design technique that divides a large table into smaller, more manageable pieces called partitions
- Each partition is a subset of the table's data, and queries can be optimized to only scan the relevant partitions, thus improving performance

Types of Partitioning

1. **Range Partitioning:** Data is divided based on a range of values
2. **List Partitioning:** Data is divided based on a list of values
3. **Hash Partitioning:** Data is divided based on a hash function
4. **Key Partitioning:** Similar to hash partitioning but uses the table's primary key

Example Scenario

- Consider a large `sales` table that contains millions of rows
- To optimize query performance, we can partition this table by the year of the `sale_date` column

01 Range Partitioning Example

1. Step 1: Create the Partitioned Table
 - We'll create a `sales` table partitioned by year

```
CREATE TABLE sales (  
  sale_id INT,  
  sale_date DATE,  
  amount DECIMAL(10, 2)
```

```

) PARTITION BY RANGE (YEAR(sale_date)) (
    PARTITION p2021 VALUES LESS THAN (2022) COMMENT = 'Data for the year 2021',
    PARTITION p2022 VALUES LESS THAN (2023) COMMENT = 'Data for the year 2022',
    PARTITION p2023 VALUES LESS THAN (2024) COMMENT = 'Data for the year 2023'
);

```

- **PARTITION BY RANGE (YEAR(sale_date))**: Partitions the table by the year extracted from the `sale_date` column
- **PARTITION p2021 VALUES LESS THAN (2022)**: All sales with a year less than 2022 are stored in the `p2021` partition
- **PARTITION p2022 VALUES LESS THAN (2023)**: All sales with a year less than 2023 are stored in the `p2022` partition
- **PARTITION p2023 VALUES LESS THAN (2024)**: All sales with a year less than 2024 are stored in the `p2023` partition

2. Step 2: Insert Data into the Partitioned Table

3. Step 3: Query the Partitioned Table

- When querying the partitioned table, the database will only scan the relevant partitions, improving performance

```

SELECT * FROM sales WHERE sale_date BETWEEN '2022-01-01' AND '2022-12-31';

```

- In this example, the database will only scan the `p2022` partition because the query is filtering by the year 2022

```

-- Viewing data in partition p2021
SELECT * FROM sales WHERE YEAR(sale_date) < 2022;

```

```

-- Viewing data in partition p2022
SELECT * FROM sales WHERE YEAR(sale_date) >= 2022 AND YEAR(sale_date) < 2023;

```

```

-- Viewing data in partition p2023
SELECT * FROM sales WHERE YEAR(sale_date) >= 2023 AND YEAR(sale_date) < 2024;

```

```
mysql> SELECT * FROM sales WHERE YEAR(sale_date) < 2022;
+-----+-----+-----+
| sale_id | sale_date | amount |
+-----+-----+-----+
|      1 | 2021-01-15 | 100.00 |
|      2 | 2021-05-20 | 150.00 |
|      3 | 2021-07-22 | 200.00 |
+-----+-----+-----+
3 rows in set (0.00 sec)

mysql> SELECT * FROM sales WHERE YEAR(sale_date) >= 2022 AND YEAR(sale_date) < 2023;
+-----+-----+-----+
| sale_id | sale_date | amount |
+-----+-----+-----+
|      4 | 2022-01-10 | 250.00 |
|      5 | 2022-03-15 | 300.00 |
|      6 | 2022-08-19 | 350.00 |
+-----+-----+-----+
3 rows in set (0.00 sec)

mysql> SELECT * FROM sales WHERE YEAR(sale_date) >= 2023 AND YEAR(sale_date) < 2024;
+-----+-----+-----+
| sale_id | sale_date | amount |
+-----+-----+-----+
|      7 | 2023-01-25 | 400.00 |
|      8 | 2023-04-05 | 450.00 |
|      9 | 2023-07-11 | 500.00 |
|     10 | 2023-10-20 | 550.00 |
+-----+-----+-----+
4 rows in set (0.01 sec)
```

Viewing Partition Information

1. Using `SHOW TABLE STATUS`

- This command provides general information about the `sales` table, including whether it is partitioned

```
SHOW TABLE STATUS LIKE 'sales';
```

```
mysql> SHOW TABLE STATUS LIKE 'sales';
+-----+-----+-----+-----+-----+-----+-----+-----+
| Name | Engine | Version | Row_format | Rows | Avg_row_length | Data_length | Max_data_l |
| te_time | Check_time | Collation | Checksum | Create_options | Comment |
+-----+-----+-----+-----+-----+-----+-----+-----+
| sales | InnoDB | 10 | Dynamic | 3 | 16384 | 49152 |
|-09-07 21:17:53 | NULL | utf8mb4_0900_ai_ci | NULL | partitioned |
+-----+-----+-----+-----+-----+-----+-----+-----+
```

2. Using `INFORMATION_SCHEMA.PARTITIONS`

- To get detailed information about the partitions, you can query the `INFORMATION_SCHEMA.PARTITIONS` table

- This query will give you information about each partition in the `sales` table

```
SELECT
    PARTITION_NAME,
    TABLE_ROWS,
    DATA_LENGTH,
    INDEX_LENGTH,
    PARTITION_COMMENT
FROM
    INFORMATION_SCHEMA.PARTITIONS
WHERE
    TABLE_SCHEMA = 'DEMO' AND
    TABLE_NAME = 'sales';
```

```
+-----+-----+-----+-----+-----+
| PARTITION_NAME | TABLE_ROWS | DATA_LENGTH | INDEX_LENGTH | PARTITION_COMMENT |
+-----+-----+-----+-----+-----+
| p2021          | 3           | 16384        | 0            | Data for the year 2021 |
| p2022          | 3           | 16384        | 0            | Data for the year 2022 |
| p2023          | 4           | 16384        | 0            | Data for the year 2023 |
+-----+-----+-----+-----+-----+
3 rows in set (0.01 sec)
```

02 List Partitioning Example

- MySQL expects integer values for the list partition
- This is because MySQL does not support list partitioning with non-integer types directly

1. Step 1: Create the Partitioned Table

- We'll create a `sales` table partitioned by the region

```
CREATE TABLE sales (
    sale_id INT,
    sale_date DATE,
    amount DECIMAL(10, 2),
    region VARCHAR(50),
    region_id INT
) PARTITION BY LIST (region_id) (
    PARTITION p_north VALUES IN (1),
    PARTITION p_south VALUES IN (2),
    PARTITION p_east VALUES IN (3),
    PARTITION p_west VALUES IN (4)
);
```

- **PARTITION BY LIST (region):** Partitions the table by the `region` column
- **PARTITION p_north VALUES IN ('North'):** All sales in the 'North' region are stored in

the `p_north` partition

- **PARTITION `p_south` VALUES IN ('South')**: All sales in the 'South' region are stored in the `p_south` partition
- **PARTITION `p_east` VALUES IN ('East')**: All sales in the 'East' region are stored in the `p_east` partition
- **PARTITION `p_west` VALUES IN ('West')**: All sales in the 'West' region are stored in the `p_west` partition

2. Step 2: Insert Data into the Partitioned Table

3. Step 3: Query the Partitioned Table

- When querying the partitioned table, the database will only scan the relevant partitions, improving performance

```
SELECT
    PARTITION_NAME,
    TABLE_ROWS,
    PARTITION_COMMENT
FROM
    INFORMATION_SCHEMA.PARTITIONS
WHERE
    TABLE_SCHEMA = 'DEMO' AND
    TABLE_NAME = 'sales';
```

PARTITION_NAME	TABLE_ROWS	PARTITION_COMMENT
p_east	2	East Region
p_north	3	North Region
p_south	3	South Region
p_west	2	West Region

4 rows in set (0.00 sec)

03 Hash Partitioning Example

1. Step 1: Create the Partitioned Table

- We'll create a `sales` table partitioned by hashing the `sale_id`

```
CREATE TABLE sales (
    sale_id INT,
    sale_date DATE,
    amount DECIMAL(10, 2)
) PARTITION BY HASH (sale_id) PARTITIONS 4;
```

- **PARTITION BY HASH (`sale_id`)**: Partitions the table by hashing the `sale_id` column

- **PARTITIONS 4:** Creates 4 partitions, with the rows distributed among them based on the hash value of `sale_id`

2. Step 2: Insert Data into the Partitioned Table

3. Step 3: Query the Partitioned Table

```
SELECT * FROM sales WHERE sale_id = 2;
```

- In this example, the database will calculate the hash value of `sale_id = 2` and scan only the relevant partition

```
SELECT
    PARTITION_NAME,
    TABLE_ROWS
FROM
    INFORMATION_SCHEMA.PARTITIONS
WHERE
    TABLE_SCHEMA = 'DEMO' AND
    TABLE_NAME = 'sales';
```

PARTITION_NAME	TABLE_ROWS
p0	2
p1	3
p2	3
p3	2

4 rows in set (0.00 sec)

Benefits of Partitioning

- **Improved Query Performance:** Queries can be optimized to scan only relevant partitions
- **Easier Maintenance:** Smaller partitions are easier to manage, back up, and restore
- **Enhanced Manageability:** Maintenance operations like archiving and purging old data can be performed on individual partitions without affecting the entire table

Subpartitioning

- In MySQL, you cannot directly apply two different types of partitioning (such as range and list partitioning) on a single table
- However, you can achieve similar functionality using subpartitioning, which allows you to first partition your table by one method (e.g., range) and then subpartition each

partition by another method (e.g., list)

Supported Partitioning and Subpartitioning Types in MySQL

1. **Range Partitioning:** Divides data into partitions based on a range of values
 - **Subpartitioning:** You can use `HASH` or `KEY` partitioning on subpartitions
2. **List Partitioning:** Divides data based on a list of discrete values
 - **Subpartitioning:** You can use `HASH` or `KEY` partitioning on subpartitions

4. Handling Large Datasets

- Handling large datasets involves techniques to manage and query vast amounts of data efficiently
- This includes sharding, indexing, and using big data technologies

Example

- Consider a `log_entries` table with billions of rows
- **Sharding:** Distribute the data across multiple database servers
- **Indexing:** Create indexes on frequently queried columns

```
CREATE INDEX idx_log_date ON log_entries (log_date);
```

- **Big Data Technologies:** Use technologies like Hadoop or Spark for distributed data processing

5. NoSQL Integration

- NoSQL databases are designed to handle unstructured data, providing flexibility and scalability
- Integrating NoSQL with relational databases can leverage the strengths of both

Example

- Consider an e-commerce application with product catalog data in a relational database and user activity logs in a NoSQL database

Relational Database

- Products Table

product_id	product_name	price
1	Laptop	1000.00
2	Smartphone	500.00

- **NoSQL Database:** User Activity Collection (MongoDB)

```
[
  {
    "user_id": "u123",
    "activity": "view",
    "product_id": "1",
    "timestamp": "2023-01-01T10:00:00Z"
  },
  {
    "user_id": "u124",
    "activity": "purchase",
    "product_id": "2",
    "timestamp": "2023-01-02T11:00:00Z"
  }
]
```

- We will learn about **NoSQL - MongoDB** deeply in **DBMS-NoSQL-MongoDB-Solutions**